

dog_app

September 13, 2019

1 Convolutional Neural Networks

1.1 Project: Write an Algorithm for a Dog Identification App

In this notebook, some template code has already been provided for you, and you will need to implement additional functionality to successfully complete this project. You will not need to modify the included code beyond what is requested. Sections that begin with '**(IMPLEMENTATION)**' in the header indicate that the following block of code will require additional functionality which you must provide. Instructions will be provided for each section, and the specifics of the implementation are marked in the code block with a 'TODO' statement. Please be sure to read the instructions carefully!

Note: Once you have completed all of the code implementations, you need to finalize your work by exporting the Jupyter Notebook as an HTML document. Before exporting the notebook to html, all of the code cells need to have been run so that reviewers can see the final implementation and output. You can then export the notebook by using the menu above and navigating to **File -> Download as -> HTML (.html)**. Include the finished document along with this notebook as your submission.

In addition to implementing code, there will be questions that you must answer which relate to the project and your implementation. Each section where you will answer a question is preceded by a '**Question X**' header. Carefully read each question and provide thorough answers in the following text boxes that begin with '**Answer:**'. Your project submission will be evaluated based on your answers to each of the questions and the implementation you provide.

Note: Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. Markdown cells can be edited by double-clicking the cell to enter edit mode.

The rubric contains *optional* "Stand Out Suggestions" for enhancing the project beyond the minimum requirements. If you decide to pursue the "Stand Out Suggestions", you should include the code in this Jupyter notebook.

Step 0: Import Datasets

Make sure that you've downloaded the required human and dog datasets:

Note: if you are using the Udacity workspace, you DO NOT need to re-download these - they can be found in the /data folder as noted in the cell below.

- Download the [dog dataset](#). Unzip the folder and place it in this project's home directory, at the location /dog_images.
- Download the [human dataset](#). Unzip the folder and place it in the home directory, at location /lfw.

Note: If you are using a Windows machine, you are encouraged to use [7zip](#) to extract the folder.

In the code cell below, we save the file paths for both the human (LFW) dataset and dog dataset in the numpy arrays human_files and dog_files.

```
In [20]: import numpy as np
         from glob import glob

         # load filenames for human and dog images
         human_files = np.array(glob("/data/lfw/*/"))
         dog_files = np.array(glob("/data/dog_images/*/"))

         # print number of images in each dataset
         print(f'There are {len(human_files)} total human images.')
         print(f'There are {len(dog_files)} total dog images.')
```

There are 13233 total human images.

There are 8351 total dog images.

Step 1: Detect Humans

In this section, we use OpenCV's implementation of [Haar feature-based cascade classifiers](#) to detect human faces in images.

OpenCV provides many pre-trained face detectors, stored as XML files on [github](#). We have downloaded one of these detectors and stored it in the haarcascades directory. In the next code cell, we demonstrate how to use this detector to find human faces in a sample image.

```
In [21]: import cv2
         import matplotlib.pyplot as plt
         %matplotlib inline

         # extract pre-trained face detector
         face_cascade = cv2.CascadeClassifier('haarcascades/haarcascade_frontalface_alt.xml')

         # load color (BGR) image
         img = cv2.imread(human_files[0])
         # convert BGR image to grayscale
         gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

         # find faces in image
         faces = face_cascade.detectMultiScale(gray)

         # print number of faces detected in the image
         print('Number of faces detected:', len(faces))
```

```

# get bounding box for each detected face
for (x,y,w,h) in faces:
    # add bounding box to color image
    cv2.rectangle(img,(x,y),(x+w,y+h),(255,0,0),2)

# convert BGR image to RGB for plotting
cv_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

# display the image, along with bounding box
plt.imshow(cv_rgb)
plt.show()

```

Number of faces detected: 1



Before using any of the face detectors, it is standard procedure to convert the images to grayscale. The `detectMultiScale` function executes the classifier stored in `face_cascade` and takes the grayscale image as a parameter.

In the above code, `faces` is a numpy array of detected faces, where each row corresponds to a detected face. Each detected face is a 1D array with four entries that specifies the bounding box of the detected face. The first two entries in the array (extracted in the above code as `x` and `y`) specify the horizontal and vertical positions of the top left corner of the bounding box. The last two entries in the array (extracted here as `w` and `h`) specify the width and height of the box.

1.1.1 Write a Human Face Detector

We can use this procedure to write a function that returns True if a human face is detected in an image and False otherwise. This function, aptly named `face_detector`, takes a string-valued file path to an image as input and appears in the code block below.

```
In [22]: # returns "True" if face is detected in image stored at img_path
def face_detector(img_path):
    img = cv2.imread(img_path)
    gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    human_faces = face_cascade.detectMultiScale(gray)
    return len(human_faces) > 0
```

1.1.2 (IMPLEMENTATION) Assess the Human Face Detector

Question 1: Use the code cell below to test the performance of the `face_detector` function.

- What percentage of the first 100 images in `human_files` have a detected human face?
- What percentage of the first 100 images in `dog_files` have a detected human face?

Ideally, we would like 100% of human images with a detected face and 0% of dog images with a detected face. You will see that our algorithm falls short of this goal, but still gives acceptable performance. We extract the file paths for the first 100 images from each of the datasets and store them in the numpy arrays `human_files_short` and `dog_files_short`.

Answer: (You can print out your results and/or write your percentages in this cell)

98 % in `human_files`

17 % in `dog_files`

```
In [23]: from tqdm import tqdm

human_files_short = human_files[:100]
dog_files_short = dog_files[:100]

#-#-# Do NOT modify the code above this line. #-#-#

## TODO: Test the performance of the face_detector algorithm
## on the images in human_files_short and dog_files_short.
list_human = [face_detector(human_files_short[i]) for i in range(0, 100)]
print (sum (list_human))

list_dog = [face_detector(dog_files_short[i]) for i in range(0, 100)]
print (sum (list_dog))
```

98

17

We suggest the face detector from OpenCV as a potential way to detect human images in your algorithm, but you are free to explore other approaches, especially approaches that make use of deep learning :). Please use the code cell below to design and test your own face detection algorithm. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

```
In [24]: ### (Optional)
        ### TODO: Test performance of another face detection algorithm.
        ### Feel free to use as many code cells as needed.
```

Step 2: Detect Dogs

In this section, we use a [pre-trained model](#) to detect dogs in images.

1.1.3 Obtain Pre-trained VGG-16 Model

The code cell below downloads the VGG-16 model, along with weights that have been trained on [ImageNet](#), a very large, very popular dataset used for image classification and other vision tasks. ImageNet contains over 10 million URLs, each linking to an image containing an object from one of 1000 categories.

```
In [25]: import torch
        import torchvision.models as models

        # define VGG16 model
        VGG16 = models.vgg16(pretrained=True)

        # check if CUDA is available
        use_cuda = torch.cuda.is_available()

        # move model to GPU if CUDA is available
        if use_cuda:
            VGG16 = VGG16.cuda()
```

Given an image, this pre-trained VGG-16 model returns a prediction (derived from the 1000 possible categories in ImageNet) for the object that is contained in the image.

1.1.4 (IMPLEMENTATION) Making Predictions with a Pre-trained Model

In the next code cell, you will write a function that accepts a path to an image (such as 'dogImages/train/001.Affenpinscher/Affenpinscher_00001.jpg') as input and returns the index corresponding to the ImageNet class that is predicted by the pre-trained VGG-16 model. The output should always be an integer between 0 and 999, inclusive.

Before writing the function, make sure that you take the time to learn how to appropriately pre-process tensors for pre-trained models in the [PyTorch documentation](#).

```
In [26]: from PIL import Image
        import torchvision.transforms as transforms
        import torchvision
        from torchvision import datasets, models, transforms
        def VGG16_predict(img_path):
            '''
            Use pre-trained VGG-16 model to obtain index corresponding to
            predicted ImageNet class for image at specified path
```

```

Args:
    img_path: path to an image

Returns:
    Index corresponding to VGG-16 model's prediction
    '''

## TODO: Complete the function.
## Load and pre-process an image from the given img_path

image = Image.open(img_path)

data_transform = transforms.Compose([transforms.RandomResizedCrop(224),
                                     transforms.ToTensor()])

image_tensor = data_transform(image)
image_tensor = image_tensor.unsqueeze(0)

# move tensor to cuda
if torch.cuda.is_available():
    image_tensor = image_tensor.cuda()

## Return the *index* of the predicted class for that image
output = VGG16(image_tensor)

if torch.cuda.is_available():
    output = output.cpu()

class_index = output.data.numpy().argmax()

return class_index # predicted class index

```

1.1.5 (IMPLEMENTATION) Write a Dog Detector

While looking at the [dictionary](#), you will notice that the categories corresponding to dogs appear in an uninterrupted sequence and correspond to dictionary keys 151-268, inclusive, to include all categories from 'Chihuahua' to 'Mexican hairless'. Thus, in order to check to see if an image is predicted to contain a dog by the pre-trained VGG-16 model, we need only check if the pre-trained model predicts an index between 151 and 268 (inclusive).

Use these ideas to complete the `dog_detector` function below, which returns `True` if a dog is detected in an image (and `False` if not).

```

In [27]: ### returns "True" if a dog is detected in the image stored at img_path
def dog_detector(img_path):

```

```

    ## TODO: Complete the function.
    class_index = VGG16_predict(img_path)
    return (151 <= class_index <= 268)

```

1.1.6 (IMPLEMENTATION) Assess the Dog Detector

Question 2: Use the code cell below to test the performance of your dog_detector function.

- What percentage of the images in human_files_short have a detected dog?
- What percentage of the images in dog_files_short have a detected dog?

Answer: 0 percentage of the images in human_files_short have a detected dog
 74 percentage of the images in dog_files_short have a detected dog

```

In [28]: ### TODO: Test the performance of the dog_detector function
         ### on the images in human_files_short and dog_files_short.
         list_human = [dog_detector(human_files_short[i]) for i in range (0, 100)]
         print (sum (list_human))

         list_dog = [dog_detector(dog_files_short[i]) for i in range (0, 100)]
         print (sum (list_dog))

```

0
 74

We suggest VGG-16 as a potential network to detect dog images in your algorithm, but you are free to explore other pre-trained networks (such as [Inception-v3](#), [ResNet-50](#), etc). Please use the code cell below to test other pre-trained PyTorch models. If you decide to pursue this *optional* task, report performance on human_files_short and dog_files_short.

```

In [29]: ### (Optional)
         ### TODO: Report the performance of another pre-trained network.
         ### Feel free to use as many code cells as needed.

```

Step 3: Create a CNN to Classify Dog Breeds (from Scratch)

Now that we have functions for detecting humans and dogs in images, we need a way to predict breed from images. In this step, you will create a CNN that classifies dog breeds. You must create your CNN *from scratch* (so, you can't use transfer learning *yet!*), and you must attain a test accuracy of at least 10%. In Step 4 of this notebook, you will have the opportunity to use transfer learning to create a CNN that attains greatly improved accuracy.

We mention that the task of assigning breed to dogs from images is considered exceptionally challenging. To see why, consider that *even a human* would have trouble distinguishing between a Brittany and a Welsh Springer Spaniel.

Brittany	Welsh Springer Spaniel
----------	------------------------

It is not difficult to find other dog breed pairs with minimal inter-class variation (for instance, Curly-Coated Retrievers and American Water Spaniels).

Curly-Coated Retriever	American Water Spaniel
------------------------	------------------------

Likewise, recall that labradors come in yellow, chocolate, and black. Your vision-based algorithm will have to conquer this high intra-class variation to determine how to classify all of these different shades as the same breed.

Yellow Labrador	Chocolate Labrador
-----------------	--------------------

We also mention that random chance presents an exceptionally low bar: setting aside the fact that the classes are slightly imbalanced, a random guess will provide a correct answer roughly 1 in 133 times, which corresponds to an accuracy of less than 1%.

Remember that the practice is far ahead of the theory in deep learning. Experiment with many different architectures, and trust your intuition. And, of course, have fun!

1.1.7 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at `dog_images/train`, `dog_images/valid`, and `dog_images/test`, respectively). You may find [this documentation on custom datasets](#) to be a useful resource. If you are interested in augmenting your training and/or validation data, check out the wide variety of [transforms](#)!

```
In [30]: import os
import torch
from torchvision import datasets
from PIL import ImageFile
ImageFile.LOAD_TRUNCATED_IMAGES = True
import torchvision.transforms as transforms

### TODO: Write data loaders for training, validation, and test sets
## Specify appropriate transforms, and batch_sizes

# define dataloader parameters
batch_size = 20
num_workers=0

# define training and test data directories
data_dir = '/data/dog_images/'
train_dir = os.path.join(data_dir, 'train/')
test_dir = os.path.join(data_dir, 'test/')
validation_dir = os.path.join(data_dir, 'valid/')
```



```

data_transform = transforms.Compose([transforms.RandomResizedCrop(224),
transforms.RandomHorizontalFlip(), # randomly flip and rotate
transforms.RandomRotation(10),
transforms.ToTensor(),
transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])

# prepare data loaders
train_data = datasets.ImageFolder(train_dir, transform=data_transform)
validation_data = datasets.ImageFolder(validation_dir, transform=data_transform)
test_data = datasets.ImageFolder(test_dir, transform=data_transform)

train_loader = torch.utils.data.DataLoader(train_data, batch_size=batch_size,
                                           num_workers=num_workers, shuffle=True)
valid_loader = torch.utils.data.DataLoader(validation_data, batch_size=batch_size,
                                           num_workers=num_workers, shuffle=True)
test_loader = torch.utils.data.DataLoader(test_data, batch_size=batch_size,
                                           num_workers=num_workers, shuffle=True)

loaders_scratch = {
    'train': train_loader,
    'valid': valid_loader,
    'test': test_loader
}

```

Question 3: Describe your chosen procedure for preprocessing the data. - How does your code resize the images (by cropping, stretching, etc)? What size did you pick for the input tensor, and why? - Did you decide to augment the dataset? If so, how (through translations, flips, rotations, etc)? If not, why not?

```

In [31]: import matplotlib.pyplot as plt
         %matplotlib inline
         import numpy as np
         # helper function to un-normalize and display an image
         def imshow(img):
             img = img / 2 + 0.5 # unnormalize
             plt.imshow(np.transpose(img, (1, 2, 0))) # convert from Tensor image

In [32]: # obtain one batch of training images
         dataiter = iter(train_loader)
         images, labels = dataiter.next()
         images = images.numpy() # convert images to numpy for display

         # plot the images in the batch, along with the corresponding labels
         fig = plt.figure(figsize=(25, 4))
         # display 20 images
         for idx in np.arange(20):

```

```
ax = fig.add_subplot(2, 20/2, idx+1, xticks=[], yticks=[])
imshow(images[idx])
```



Answer: Images are resized using RandomResizedCrop. Size 224 was picked for input tensor by looking at the published articles online. That seemed the most appropriate size to take. Yes, the dataset was augmented by the techniques mentioned in the videos using RandomHorizontalFlip, RandomRotation. Image file was finally converted to tensor and normalized.

1.1.8 (IMPLEMENTATION) Model Architecture

Create a CNN to classify dog breed. Use the template in the code cell below.

```
In [33]: import torch.nn as nn
import torch.nn.functional as F

# define the CNN architecture
class Net(nn.Module):
    ### TODO: choose an architecture, and complete the class
    def __init__(self):
        super(Net, self).__init__()
        ## Define layers of a CNN
        # convolutional layer (sees 224x224x3 image tensor)
        self.conv1 = nn.Conv2d(3, 32, 3, padding=1)
        # convolutional layer (sees 112x112x16 tensor)
        self.conv2 = nn.Conv2d(32, 64, 3, padding=1)
        # convolutional layer (sees 64x64x32 tensor)
        self.conv3 = nn.Conv2d(64, 128, 3, padding=1)
        # max pooling layer
        self.pool = nn.MaxPool2d(2, 2)
        # linear layer (128 * 28 * 28 -> 500)
        self.fc1 = nn.Linear(128 * 28 * 28, 500)
        # linear layer (500 -> 133)
        self.fc2 = nn.Linear(500, 133)
        # dropout layer (p=0.25)
        self.dropout = nn.Dropout(0.25)
    def forward(self, x):
```

```

    ## Define forward behavior
    x = self.pool(F.relu(self.conv1(x)))
    x = self.pool(F.relu(self.conv2(x)))
    x = self.pool(F.relu(self.conv3(x)))
    # flatten image input
    x = x.view(-1, 128 * 28 * 28)
    # add dropout layer
    x = self.dropout(x)
    # add 1st hidden layer, with relu activation function
    x = F.relu(self.fc1(x))
    # add dropout layer
    x = self.dropout(x)
    # add 2nd hidden layer, with relu activation function
    x = self.fc2(x)
    return x

##-## You so NOT have to modify the code below this line. ##-##

# instantiate the CNN
model_scratch = Net()

# move tensors to GPU if CUDA is available
if use_cuda:
    model_scratch.cuda()

```

Question 4: Outline the steps you took to get to your final CNN architecture and your reasoning at each step.

Answer: 3 Convolutional layer and 2 linear layer CNN was designed. Relu activation function used with pooling of 2x. 1st layer gets input image of 224 224 3 dimensions and has kernel size of (3,3), and 32 filtered images as output. ReLu activation function is used and pooling would downsample dim to 112 112 3. 2nd layer would take 32 images, with output of 64. Relu function is used and pooling would downsize dim to 56 56 4. 3rd layer would take 64 as input and 128 as output, and relu function and pooling would downsize dimension to 28 28 128. This would be input for fully connected layer 1 with output as 500 classes. 2nd fully connected layer would take these 500 as inputs and would give 133 classes as output using relu function.

1.1.9 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_scratch`, and the optimizer as `optimizer_scratch` below.

```

In [34]: import torch.optim as optim

### TODO: select loss function
criterion_scratch = nn.CrossEntropyLoss()

### TODO: select optimizer
optimizer_scratch = optim.SGD(model_scratch.parameters(), lr=0.01)

```

```

if use_cuda:
    criterion_scratch = criterion_scratch.cuda()

```

1.1.10 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath 'model_scratch.pt'.

```

In [35]: def train(n_epochs, loaders, model, optimizer, criterion, use_cuda, save_path):
    """returns trained model"""
    # initialize tracker for minimum validation loss
    valid_loss_min = np.Inf

    for epoch in range(1, n_epochs+1):
        # initialize loss variables for training loss and validation loss
        train_loss = 0.0
        valid_loss = 0.0

        # train the model

        model.train()
        for batch_idx, (data, target) in enumerate(loaders['train']):
            # move to GPU
            if use_cuda:
                data, target = data.cuda(), target.cuda()
            ## update the model parameters
            optimizer.zero_grad()
            output = model(data)
            loss = criterion(output, target)
            loss.backward()
            optimizer.step()

            ## average training loss
            train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))

        # validate the model
        model.eval()
        for batch_idx, (data, target) in enumerate(loaders['valid']):
            # move to GPU
            if use_cuda:
                data, target = data.cuda(), target.cuda()
            output = model(data)
            # calculate loss and update average validation loss
            loss = criterion(output, target)
            valid_loss = valid_loss + ((1 / (batch_idx + 1)) * (loss.data - valid_loss))

```

```

# training/validation loss
print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(
    epoch,
    train_loss,
    valid_loss
))

## save the model if validation loss has decreased
if valid_loss <= valid_loss_min:
    print(f'Validation loss decreased ({valid_loss_min} --> {valid_loss}). Saving model ...')
    torch.save(model.state_dict(), save_path)
    valid_loss_min = valid_loss

# return trained model
return model

# train the model
model_scratch = train(50, loaders_scratch, model_scratch, optimizer_scratch, criterion_scratch)

# load the model with best validation accuracy
model_scratch.load_state_dict(torch.load('model_scratch.pt'))

```

```

Epoch: 1      Training Loss: 4.883908      Validation Loss: 4.867745
Validation loss decreased (inf --> 4.8677449226379395). Saving model ...
Epoch: 2      Training Loss: 4.857948      Validation Loss: 4.828182
Validation loss decreased (4.8677449226379395 --> 4.828181743621826). Saving model ...
Epoch: 3      Training Loss: 4.776867      Validation Loss: 4.714515
Validation loss decreased (4.828181743621826 --> 4.714515209197998). Saving model ...
Epoch: 4      Training Loss: 4.653889      Validation Loss: 4.627423
Validation loss decreased (4.714515209197998 --> 4.62742280960083). Saving model ...
Epoch: 5      Training Loss: 4.603976      Validation Loss: 4.595919
Validation loss decreased (4.62742280960083 --> 4.595919132232666). Saving model ...
Epoch: 6      Training Loss: 4.562114      Validation Loss: 4.592135
Validation loss decreased (4.595919132232666 --> 4.592134952545166). Saving model ...
Epoch: 7      Training Loss: 4.527866      Validation Loss: 4.517825
Validation loss decreased (4.592134952545166 --> 4.517824649810791). Saving model ...
Epoch: 8      Training Loss: 4.497367      Validation Loss: 4.512098
Validation loss decreased (4.517824649810791 --> 4.51209831237793). Saving model ...
Epoch: 9      Training Loss: 4.473951      Validation Loss: 4.484348
Validation loss decreased (4.51209831237793 --> 4.484348297119141). Saving model ...
Epoch: 10     Training Loss: 4.437525      Validation Loss: 4.456133
Validation loss decreased (4.484348297119141 --> 4.456132888793945). Saving model ...
Epoch: 11     Training Loss: 4.391824      Validation Loss: 4.421599
Validation loss decreased (4.456132888793945 --> 4.4215989112854). Saving model ...
Epoch: 12     Training Loss: 4.370867      Validation Loss: 4.465789
Epoch: 13     Training Loss: 4.345278      Validation Loss: 4.371666
Validation loss decreased (4.4215989112854 --> 4.371665954589844). Saving model ...
Epoch: 14     Training Loss: 4.327038      Validation Loss: 4.431526

```

Epoch: 15	Training Loss: 4.314827	Validation Loss: 4.408340
Epoch: 16	Training Loss: 4.274291	Validation Loss: 4.339738
Validation loss decreased (4.371665954589844 --> 4.339737892150879). Saving model ...		
Epoch: 17	Training Loss: 4.255634	Validation Loss: 4.365775
Epoch: 18	Training Loss: 4.236125	Validation Loss: 4.359012
Epoch: 19	Training Loss: 4.199792	Validation Loss: 4.341966
Epoch: 20	Training Loss: 4.168004	Validation Loss: 4.377005
Epoch: 21	Training Loss: 4.148744	Validation Loss: 4.310865
Validation loss decreased (4.339737892150879 --> 4.3108649253845215). Saving model ...		
Epoch: 22	Training Loss: 4.126297	Validation Loss: 4.322583
Epoch: 23	Training Loss: 4.105670	Validation Loss: 4.290994
Validation loss decreased (4.3108649253845215 --> 4.290994167327881). Saving model ...		
Epoch: 24	Training Loss: 4.078447	Validation Loss: 4.200477
Validation loss decreased (4.290994167327881 --> 4.200477123260498). Saving model ...		
Epoch: 25	Training Loss: 4.049450	Validation Loss: 4.226094
Epoch: 26	Training Loss: 4.030415	Validation Loss: 4.183950
Validation loss decreased (4.200477123260498 --> 4.183949947357178). Saving model ...		
Epoch: 27	Training Loss: 4.002006	Validation Loss: 4.247792
Epoch: 28	Training Loss: 3.970855	Validation Loss: 4.154358
Validation loss decreased (4.183949947357178 --> 4.15435791015625). Saving model ...		
Epoch: 29	Training Loss: 3.944093	Validation Loss: 4.135512
Validation loss decreased (4.15435791015625 --> 4.135512351989746). Saving model ...		
Epoch: 30	Training Loss: 3.925878	Validation Loss: 4.113176
Validation loss decreased (4.135512351989746 --> 4.113176345825195). Saving model ...		
Epoch: 31	Training Loss: 3.911366	Validation Loss: 4.118761
Epoch: 32	Training Loss: 3.874173	Validation Loss: 4.112432
Validation loss decreased (4.113176345825195 --> 4.112431526184082). Saving model ...		
Epoch: 33	Training Loss: 3.822076	Validation Loss: 4.083956
Validation loss decreased (4.112431526184082 --> 4.083956241607666). Saving model ...		
Epoch: 34	Training Loss: 3.807226	Validation Loss: 4.134215
Epoch: 35	Training Loss: 3.796986	Validation Loss: 4.182295
Epoch: 36	Training Loss: 3.798250	Validation Loss: 4.070650
Validation loss decreased (4.083956241607666 --> 4.07064962387085). Saving model ...		
Epoch: 37	Training Loss: 3.748355	Validation Loss: 4.115151
Epoch: 38	Training Loss: 3.737085	Validation Loss: 4.053328
Validation loss decreased (4.07064962387085 --> 4.053328037261963). Saving model ...		
Epoch: 39	Training Loss: 3.716938	Validation Loss: 4.136744
Epoch: 40	Training Loss: 3.678439	Validation Loss: 4.066415
Epoch: 41	Training Loss: 3.663081	Validation Loss: 4.055721
Epoch: 42	Training Loss: 3.651802	Validation Loss: 4.059011
Epoch: 43	Training Loss: 3.607549	Validation Loss: 4.217405
Epoch: 44	Training Loss: 3.590004	Validation Loss: 4.035468
Validation loss decreased (4.053328037261963 --> 4.035468101501465). Saving model ...		
Epoch: 45	Training Loss: 3.592422	Validation Loss: 4.007452
Validation loss decreased (4.035468101501465 --> 4.00745153427124). Saving model ...		
Epoch: 46	Training Loss: 3.565088	Validation Loss: 4.053574
Epoch: 47	Training Loss: 3.537142	Validation Loss: 4.098477
Epoch: 48	Training Loss: 3.493073	Validation Loss: 4.099118

```
Epoch: 49           Training Loss: 3.471225           Validation Loss: 4.001259
Validation loss decreased (4.00745153427124 --> 4.0012593269348145). Saving model ...
Epoch: 50           Training Loss: 3.496581           Validation Loss: 3.980536
Validation loss decreased (4.0012593269348145 --> 3.9805357456207275). Saving model ...
```

1.1.11 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 10%.

```
In [36]: def test(loaders, model, criterion, use_cuda):

    # monitor test loss and accuracy
    test_loss = 0.
    correct = 0.
    total = 0.

    model.eval()
    for batch_idx, (data, target) in enumerate(loaders['test']):
        # move to GPU
        if use_cuda:
            data, target = data.cuda(), target.cuda()
        # forward pass: compute predicted outputs by passing inputs to the model
        output = model(data)
        # calculate the loss
        loss = criterion(output, target)
        # update average test loss
        test_loss = test_loss + ((1 / (batch_idx + 1)) * (loss.data - test_loss))
        # convert output probabilities to predicted class
        pred = output.data.max(1, keepdim=True)[1]
        # compare predictions to true label
        correct += np.sum(np.squeeze(pred.eq(target.data.view_as(pred))).cpu().numpy())
        total += data.size(0)

    print(f'Test Loss: {test_loss}')

    print('\nTest Accuracy: %2d%% (%2d/%2d)' % (
        100. * correct / total, correct, total))

    # call test function
    test(loaders_scratch, model_scratch, criterion_scratch, use_cuda)
```

```
Test Loss: 3.9477438926696777
```

```
Test Accuracy: 11% (97/836)
```

Step 4: Create a CNN to Classify Dog Breeds (using Transfer Learning)

You will now use transfer learning to create a CNN that can identify dog breed from images. Your CNN must attain at least 60% accuracy on the test set.

1.1.12 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at dogImages/train, dogImages/valid, and dogImages/test, respectively).

If you like, **you are welcome to use the same data loaders from the previous step**, when you created a CNN from scratch.

```
In [37]: #Specify data loaders
        loaders_transfer = loaders_scratch.copy()
```

1.1.13 (IMPLEMENTATION) Model Architecture

Use transfer learning to create a CNN to classify dog breed. Use the code cell below, and save your initialized model as the variable `model_transfer`.

```
In [38]: import torchvision.models as models
        import torch.nn as nn

        ## TODO: Specify model architecture
        model_transfer = models.resnet50(pretrained=True)

        for param in model_transfer.parameters():
            param.requires_grad = False

        model_transfer.fc = nn.Linear(2048, 133, bias=True)

        model_fc_parameters = model_transfer.fc.parameters()

        for param in model_fc_parameters:
            param.requires_grad = True

        if use_cuda:
            model_transfer = model_transfer.cuda()
```

```
Downloading: "https://download.pytorch.org/models/resnet50-19c8e357.pth" to /root/.torch/models/
100%|| 102502400/102502400 [00:01<00:00, 75585018.79it/s]
```

Question 5: Outline the steps you took to get to your final CNN architecture and your reasoning at each step. Describe why you think the architecture is suitable for the current problem.

Answer:

ResNet is chosen as this model was the winner of ImageNet challenge in 2015. It is a good starting point for transfer learning. Final fully connected layer output is 133 as that's the total classes of dog we have.

1.1.14 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_transfer`, and the optimizer as `optimizer_transfer` below.

```
In [39]: criterion_transfer = nn.CrossEntropyLoss()
         optimizer_transfer = optim.SGD(model_transfer.fc.parameters(), lr=0.001)
```

1.1.15 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath `'model_transfer.pt'`.

```
In [40]: def train(n_epochs, loaders, model, optimizer, criterion, use_cuda, save_path):
         """returns trained model"""

         # initialize tracker for minimum validation loss
         valid_loss_min = np.Inf

         for epoch in range(1, n_epochs+1):
             # initialize loss variables for training and validation loss
             train_loss = 0.0
             valid_loss = 0.0

             # train the model
             model.train()
             for batch_idx, (data, target) in enumerate(loaders['train']):
                 # move to GPU
                 if use_cuda:
                     data, target = data.cuda(), target.cuda()

                 optimizer.zero_grad()
                 # forward pass: compute predicted outputs by passing inputs to the model
                 output = model(data)
                 # calculate the batch loss
                 loss = criterion(output, target)
                 # backward pass: compute gradient of the loss with respect to model parameters
                 loss.backward()
                 # perform a single optimization step (parameter update)
                 optimizer.step()

                 train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))

             # validate the model
             model.eval()
             for batch_idx, (data, target) in enumerate(loaders['valid']):
                 # move to GPU
                 if use_cuda:
```

```

        data, target = data.cuda(), target.cuda()
        ## update the average validation loss
        output = model(data)
        loss = criterion(output, target)
        valid_loss = valid_loss + ((1 / (batch_idx + 1)) * (loss.data - valid_loss))

    # print training/validation loss
    print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(
        epoch,
        train_loss,
        valid_loss
    ))

    ## save the model if validation loss has decreased
    if valid_loss < valid_loss_min:
        torch.save(model.state_dict(), save_path)
        print('Validation loss decreased ({:.6f} --> {:.6f}). Saving model ...'
              .format(valid_loss_min, valid_loss))
        valid_loss_min = valid_loss

    # return trained model
    return model

n_epochs = 50

# train the model
model_transfer = train(n_epochs, loaders_transfer, model_transfer, optimizer_transfer,
                       use_cuda, 'model_transfer.pt')

# load the model that got the best validation accuracy
model_transfer.load_state_dict(torch.load('model_transfer.pt'))

```

```

Epoch: 1      Training Loss: 4.825168      Validation Loss: 4.703339
Validation loss decreased (inf --> 4.703339). Saving model ...
Epoch: 2      Training Loss: 4.629292      Validation Loss: 4.499588
Validation loss decreased (4.703339 --> 4.499588). Saving model ...
Epoch: 3      Training Loss: 4.458992      Validation Loss: 4.306858
Validation loss decreased (4.499588 --> 4.306858). Saving model ...
Epoch: 4      Training Loss: 4.290055      Validation Loss: 4.156352
Validation loss decreased (4.306858 --> 4.156352). Saving model ...
Epoch: 5      Training Loss: 4.129441      Validation Loss: 3.987661
Validation loss decreased (4.156352 --> 3.987661). Saving model ...
Epoch: 6      Training Loss: 3.979967      Validation Loss: 3.838576
Validation loss decreased (3.987661 --> 3.838576). Saving model ...
Epoch: 7      Training Loss: 3.828594      Validation Loss: 3.672453
Validation loss decreased (3.838576 --> 3.672453). Saving model ...
Epoch: 8      Training Loss: 3.711260      Validation Loss: 3.463255
Validation loss decreased (3.672453 --> 3.463255). Saving model ...

```

Epoch: 9	Training Loss: 3.566027	Validation Loss: 3.395427
Validation loss decreased (3.463255 --> 3.395427). Saving model ...		
Epoch: 10	Training Loss: 3.437756	Validation Loss: 3.253336
Validation loss decreased (3.395427 --> 3.253336). Saving model ...		
Epoch: 11	Training Loss: 3.337365	Validation Loss: 3.119250
Validation loss decreased (3.253336 --> 3.119250). Saving model ...		
Epoch: 12	Training Loss: 3.217595	Validation Loss: 3.015644
Validation loss decreased (3.119250 --> 3.015644). Saving model ...		
Epoch: 13	Training Loss: 3.130946	Validation Loss: 2.875268
Validation loss decreased (3.015644 --> 2.875268). Saving model ...		
Epoch: 14	Training Loss: 3.003729	Validation Loss: 2.834037
Validation loss decreased (2.875268 --> 2.834037). Saving model ...		
Epoch: 15	Training Loss: 2.914529	Validation Loss: 2.729851
Validation loss decreased (2.834037 --> 2.729851). Saving model ...		
Epoch: 16	Training Loss: 2.849252	Validation Loss: 2.645480
Validation loss decreased (2.729851 --> 2.645480). Saving model ...		
Epoch: 17	Training Loss: 2.767242	Validation Loss: 2.594935
Validation loss decreased (2.645480 --> 2.594935). Saving model ...		
Epoch: 18	Training Loss: 2.692084	Validation Loss: 2.504748
Validation loss decreased (2.594935 --> 2.504748). Saving model ...		
Epoch: 19	Training Loss: 2.613697	Validation Loss: 2.431084
Validation loss decreased (2.504748 --> 2.431084). Saving model ...		
Epoch: 20	Training Loss: 2.542523	Validation Loss: 2.380444
Validation loss decreased (2.431084 --> 2.380444). Saving model ...		
Epoch: 21	Training Loss: 2.483855	Validation Loss: 2.281067
Validation loss decreased (2.380444 --> 2.281067). Saving model ...		
Epoch: 22	Training Loss: 2.425716	Validation Loss: 2.238971
Validation loss decreased (2.281067 --> 2.238971). Saving model ...		
Epoch: 23	Training Loss: 2.371985	Validation Loss: 2.116541
Validation loss decreased (2.238971 --> 2.116541). Saving model ...		
Epoch: 24	Training Loss: 2.325711	Validation Loss: 2.174445
Epoch: 25	Training Loss: 2.272822	Validation Loss: 2.090773
Validation loss decreased (2.116541 --> 2.090773). Saving model ...		
Epoch: 26	Training Loss: 2.222103	Validation Loss: 2.046500
Validation loss decreased (2.090773 --> 2.046500). Saving model ...		
Epoch: 27	Training Loss: 2.186184	Validation Loss: 1.997470
Validation loss decreased (2.046500 --> 1.997470). Saving model ...		
Epoch: 28	Training Loss: 2.143927	Validation Loss: 1.871257
Validation loss decreased (1.997470 --> 1.871257). Saving model ...		
Epoch: 29	Training Loss: 2.100025	Validation Loss: 1.967346
Epoch: 30	Training Loss: 2.066503	Validation Loss: 1.874444
Epoch: 31	Training Loss: 2.046357	Validation Loss: 1.835734
Validation loss decreased (1.871257 --> 1.835734). Saving model ...		
Epoch: 32	Training Loss: 2.016487	Validation Loss: 1.758526
Validation loss decreased (1.835734 --> 1.758526). Saving model ...		
Epoch: 33	Training Loss: 1.972798	Validation Loss: 1.775624
Epoch: 34	Training Loss: 1.920489	Validation Loss: 1.800687
Epoch: 35	Training Loss: 1.896557	Validation Loss: 1.716136

```

Validation loss decreased (1.758526 --> 1.716136). Saving model ...
Epoch: 36      Training Loss: 1.872892      Validation Loss: 1.726189
Epoch: 37      Training Loss: 1.849016      Validation Loss: 1.662001
Validation loss decreased (1.716136 --> 1.662001). Saving model ...
Epoch: 38      Training Loss: 1.848487      Validation Loss: 1.664742
Epoch: 39      Training Loss: 1.798081      Validation Loss: 1.596139
Validation loss decreased (1.662001 --> 1.596139). Saving model ...
Epoch: 40      Training Loss: 1.780123      Validation Loss: 1.637440
Epoch: 41      Training Loss: 1.762377      Validation Loss: 1.602790
Epoch: 42      Training Loss: 1.726585      Validation Loss: 1.523098
Validation loss decreased (1.596139 --> 1.523098). Saving model ...
Epoch: 43      Training Loss: 1.713249      Validation Loss: 1.517028
Validation loss decreased (1.523098 --> 1.517028). Saving model ...
Epoch: 44      Training Loss: 1.726695      Validation Loss: 1.531536
Epoch: 45      Training Loss: 1.684299      Validation Loss: 1.506107
Validation loss decreased (1.517028 --> 1.506107). Saving model ...
Epoch: 46      Training Loss: 1.645051      Validation Loss: 1.518529
Epoch: 47      Training Loss: 1.653809      Validation Loss: 1.527501
Epoch: 48      Training Loss: 1.630452      Validation Loss: 1.549457
Epoch: 49      Training Loss: 1.603473      Validation Loss: 1.483806
Validation loss decreased (1.506107 --> 1.483806). Saving model ...
Epoch: 50      Training Loss: 1.599444      Validation Loss: 1.427736
Validation loss decreased (1.483806 --> 1.427736). Saving model ...

```

1.1.16 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 60%.

```
In [41]: test(loaders_transfer, model_transfer, criterion_transfer, use_cuda)
```

```
Test Loss: 1.4854999780654907
```

```
Test Accuracy: 69% (583/836)
```

1.1.17 (IMPLEMENTATION) Predict Dog Breed with the Model

Write a function that takes an image path as input and returns the dog breed (Affenpinscher, Afghan hound, etc) that is predicted by your model.

```
In [42]: from PIL import Image
import torchvision.transforms as transforms

data_transfer = loaders_transfer.copy()

# list of class names by index, i.e. a name can be accessed like class_names[0]
class_names = [item[4:].replace("_", " ") for item in data_transfer['train'].dataset.cl
```

```

def predict_breed_transfer(img_path):
    ### a function that takes a path to an image as input
    ### and returns the dog breed that is predicted by the model.

    global model_transfer
    global data_transform

    # load the image and return the predicted breed
    image = Image.open(img_path).convert('RGB')

    # Removing transparent, alpha
    image = data_transform(image)[:3,:,:].unsqueeze(0)

    if use_cuda:
        model_transfer = model_transfer.cuda()
        image = image.cuda()

    model_transfer.eval()
    idx = torch.argmax(model_transfer(image))
    return class_names[idx]

```

FileNotFoundError Traceback (most recent call last)

```

<ipython-input-42-e94fff51cec6> in <module>()
    32
    33 #Test predict_breed_transfer
---> 34 for img_file in os.listdir('dogImages/test/001.Affenpinscher'):
    35     img_path = os.path.join('dogImages/test/001.Affenpinscher', img_file)
    36     predicted_breed = predict_breed_transfer(img_path)

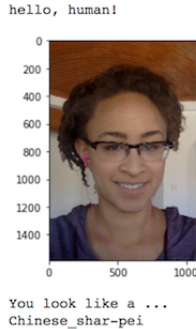
```

FileNotFoundError: [Errno 2] No such file or directory: 'dogImages/test/001.Affenpinsche

Step 5: Write your Algorithm

Write an algorithm that accepts a file path to an image and first determines whether the image contains a human, dog, or neither. Then, - if a **dog** is detected in the image, return the predicted breed. - if a **human** is detected in the image, return the resembling dog breed. - if **neither** is detected in the image, provide output that indicates an error.

You are welcome to write your own functions for detecting humans and dogs in images, but feel free to use the `face_detector` and `human_detector` functions developed above. You are **required** to use your CNN from Step 4 to predict dog breed.



Sample Human Output

Some sample output for our algorithm is provided below, but feel free to design your own user experience!

1.1.18 (IMPLEMENTATION) Write your Algorithm

```
In [43]: def run_app(img_path):
    ## handle cases for a human face, dog, and neither
    if face_detector(img_path) > 0:
        breed = predict_breed_transfer(img_path)
        print('Human / resembling dog breed is ' + breed)
    elif dog_detector(img_path):
        breed = predict_breed_transfer(img_path)
        print('Dog / predicted dog breed is ' + breed)
    else:
        print('Not Human, Neither Dog')
```

Step 6: Test Your Algorithm

In this section, you will take your new algorithm for a spin! What kind of dog does the algorithm think that *you* look like? If you have a dog, does it predict your dog's breed accurately? If you have a cat, does it mistakenly think that your cat is a dog?

1.1.19 (IMPLEMENTATION) Test Your Algorithm on Sample Images!

Test your algorithm at least six images on your computer. Feel free to use any images you like. Use at least two human and two dog images.

Question 6: Is the output better than you expected :) ? Or worse :(? Provide at least three possible points of improvement for your algorithm.

Answer:

The output is as I expected but definitely could be improved further.

Three possible points for improvement: 1. More datapoints for each dog class would improve accuracy. 2. Better transfer learning models may be more layers 3. More epochs is increasing score definitely but computationally expensive 4. Playing with different layers output

```
In [44]: ## TODO: Execute your algorithm from Step 6 on
    ## at least 6 images on your computer.
```

```


## Feel free to use as many code cells as needed.
import os

for img_file in os.listdir('human_dir/'):
    img_path = os.path.join('human_dir', img_file)
    run_app(img_path)
    print(img_path)
    img = Image.open(img_path)
    plt.imshow(img)
    plt.show()

for img_file in os.listdir('dog_dir/'):
    img_path = os.path.join('dog_dir', img_file)
    run_app(img_path)
    print(img_path)
    img = Image.open(img_path)
    plt.imshow(img)
    plt.show()

```

Human / resembling dog breed is Bullmastiff
human_dir/human3.jpg




Not Dog, Neither Human
human_dir/human7.jpeg



Human / resembing dog breed is Italian greyhound
human_dir/human1.jpg




Not Dog, Neither Human
human_dir/human4.png



Human / resembling dog breed is Greyhound
human_dir/human2.jpg



Not Dog, Neither Human
human_dir/human6.jpg



Human / resembing dog breed is Dachshund
human_dir/human8.jpeg



error Traceback (most recent call last)

```
<ipython-input-44-2879e1a9971b> in <module>()
      6 for img_file in os.listdir('human_dir/'):
      7     img_path = os.path.join('human_dir', img_file)
----> 8     run_app(img_path)
      9     print(img_path)
     10     img = Image.open(img_path)

<ipython-input-43-eae0f70da8a1> in run_app(img_path)
      3 def run_app(img_path):
      4     ## handle cases for a human face, dog, and neither
----> 5     if face_detector(img_path) > 0:
      6         breed = predict_breed_transfer(img_path)
      7         print('Human / resembling dog breed is ' + breed)

<ipython-input-22-9d98894748d5> in face_detector(img_path)
      2 def face_detector(img_path):
      3     img = cv2.imread(img_path)
----> 4     gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
      5     faces = face_cascade.detectMultiScale(gray)
      6     return len(faces) > 0
```

error: /tmp/build/80754af9/opencv_1512491964794/work/modules/imgproc/src/color.cpp:11048

In []: